

LoanApprovalPrediction

June 19, 2023

1 Loan Approval Prediction

1.0.1 Introduction:

The Loan Approval Prediction project aims to analyze a dataset related to loan applications and approvals to gain insights into the factors influencing loan approval decisions and explore the possibility of making predictions based on the analysis.

1.0.2 The primary objectives of this project are as follows:

Univariate Analysis: To examine each feature individually, studying their distributions and characteristics. This analysis will provide a comprehensive understanding of the dataset and the range of values for each feature.

Bivariate Analysis: To explore the relationships between different features and loan approval status. By examining the correlations and associations between variables, we can identify key factors that impact the loan approval decision.

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[2]: loan_appr = pd.read_csv('Loan+Approval+Prediction.csv')
print(loan_appr.head(5))
print('\n')
print(loan_appr.tail(5))
print('\n')
print(loan_appr.info())
print(loan_appr.describe())
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
0	LP001002	Male	No	0	Graduate	No	
1	LP001003	Male	Yes	1	Graduate	No	
2	LP001005	Male	Yes	0	Graduate	Yes	
3	LP001006	Male	Yes	0	Not Graduate	No	
4	LP001008	Male	No	0	Graduate	No	

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\
0	5849	0.0	NaN	360.0	

1	4583	1508.0	128.0	360.0
2	3000	0.0	66.0	360.0
3	2583	2358.0	120.0	360.0
4	6000	0.0	141.0	360.0

	Credit_History	Property_Area	Loan_Status
0	1.0	Urban	Y
1	1.0	Rural	N
2	1.0	Urban	Y
3	1.0	Urban	Y
4	1.0	Urban	Y

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
609	LP002978	Female	No	0	Graduate	No	
610	LP002979	Male	Yes	3+	Graduate	No	
611	LP002983	Male	Yes	1	Graduate	No	
612	LP002984	Male	Yes	2	Graduate	No	
613	LP002990	Female	No	0	Graduate	Yes	

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\
609	2900	0.0	71.0	360.0	
610	4106	0.0	40.0	180.0	
611	8072	240.0	253.0	360.0	
612	7583	0.0	187.0	360.0	
613	4583	0.0	133.0	360.0	

	Credit_History	Property_Area	Loan_Status
609	1.0	Rural	Y
610	1.0	Rural	Y
611	1.0	Urban	Y
612	1.0	Urban	Y
613	0.0	Semiurban	N

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 614 entries, 0 to 613
```

```
Data columns (total 13 columns):
```

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	Loan_ID	614 non-null	object
1	Gender	601 non-null	object
2	Married	611 non-null	object
3	Dependents	599 non-null	object
4	Education	614 non-null	object
5	Self_Employed	582 non-null	object
6	ApplicantIncome	614 non-null	int64
7	CoapplicantIncome	614 non-null	float64

```

8   LoanAmount          592 non-null   float64
9   Loan_Amount_Term    600 non-null   float64
10  Credit_History      564 non-null   float64
11  Property_Area       614 non-null   object
12  Loan_Status         614 non-null   object

```

```
dtypes: float64(4), int64(1), object(8)
```

```
memory usage: 62.5+ KB
```

```
None
```

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term \
count	614.000000	614.000000	592.000000	600.000000
mean	5403.459283	1621.245798	146.412162	342.000000
std	6109.041673	2926.248369	85.587325	65.12041
min	150.000000	0.000000	9.000000	12.000000
25%	2877.500000	0.000000	100.000000	360.000000
50%	3812.500000	1188.500000	128.000000	360.000000
75%	5795.000000	2297.250000	168.000000	360.000000
max	81000.000000	41667.000000	700.000000	480.000000

	Credit_History
count	564.000000
mean	0.842199
std	0.364878
min	0.000000
25%	1.000000
50%	1.000000
75%	1.000000
max	1.000000

```
[3]: loan_appr.shape
```

```
[3]: (614, 13)
```

We can observe the Dataset has 614 rows and 13 columns

```
[4]: print(loan_appr.columns)
loan_appr.dtypes
```

```
Index(['Loan_ID', 'Gender', 'Married', 'Dependents', 'Education',
       'Self_Employed', 'ApplicantIncome', 'CoapplicantIncome', 'LoanAmount',
       'Loan_Amount_Term', 'Credit_History', 'Property_Area', 'Loan_Status'],
      dtype='object')
```

```
[4]: Loan_ID          object
Gender              object
Married            object
Dependents         object
Education          object
Self_Employed     object
ApplicantIncome    int64
```

```

CoapplicantIncome    float64
LoanAmount           float64
Loan_Amount_Term     float64
Credit_History       float64
Property_Area        object
Loan_Status          object
dtype: object

```

```

[71]: print(loan_appr['Gender'].value_counts())
      print(loan_appr['Education'].value_counts())
      print(loan_appr['Property_Area'].value_counts())
      print(loan_appr['Loan_Status'].value_counts())
      print(loan_appr['Loan_Amount_Term'].value_counts())
      print(loan_appr['Credit_History'].value_counts())

```

```

Male      489
Female    112
Name: Gender, dtype: int64
Graduate      480
Not Graduate  134
Name: Education, dtype: int64
Semiurban    233
Urban        202
Rural        179
Name: Property_Area, dtype: int64
Y      422
N      192
Name: Loan_Status, dtype: int64
360.0      512
180.0       44
480.0       15
300.0       13
240.0        4
84.0         4
120.0        3
60.0         2
36.0         2
12.0         1
Name: Loan_Amount_Term, dtype: int64
1.0      475
0.0       89
Name: Credit_History, dtype: int64

```

We can start by plotting the following individual variables to get a better understanding of loan approval, and it's prediction: - 'Loan_Status': This is the target variable that indicates whether a loan was approved or not. Plotting the distribution of loan approvals ('Y') versus rejections ('N') will provide insights into the overall loan approval rate.


- 'gender': Plotting the count or proportion of loan approvals based on gender can help

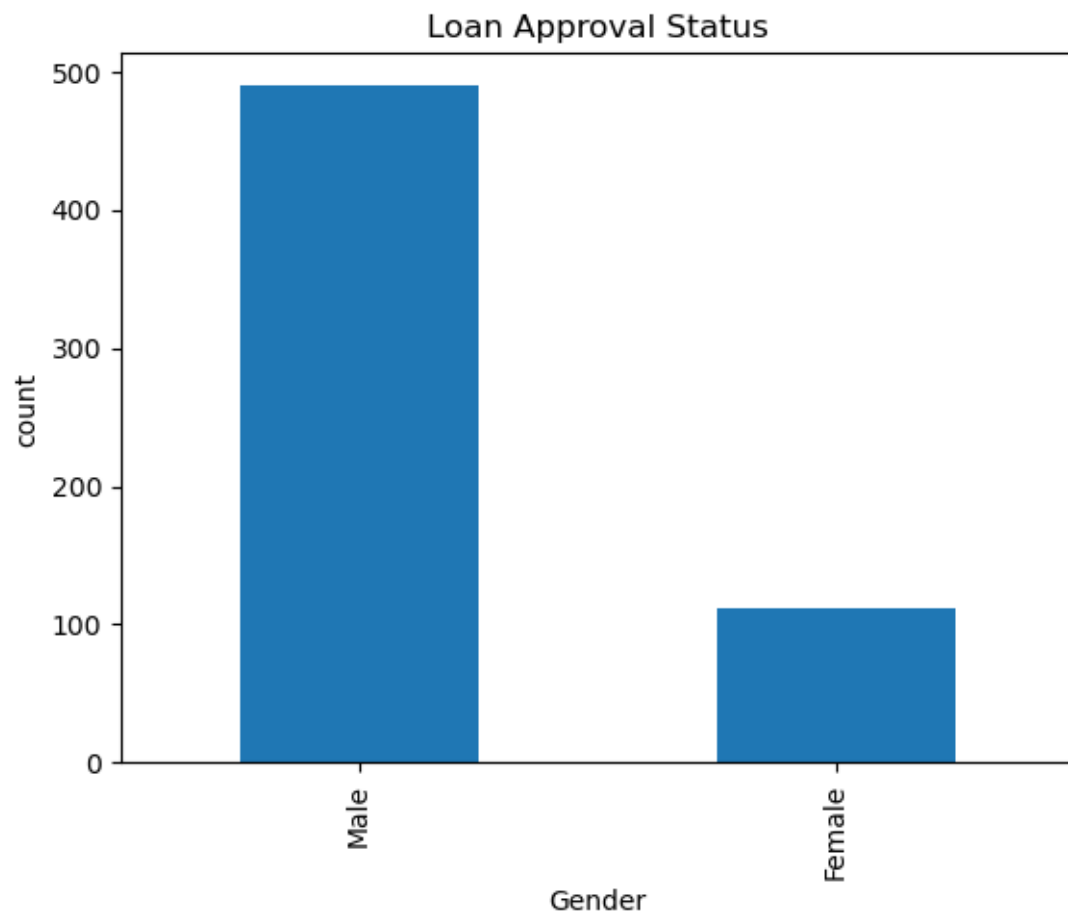
identify any gender-related patterns in loan approval rates.

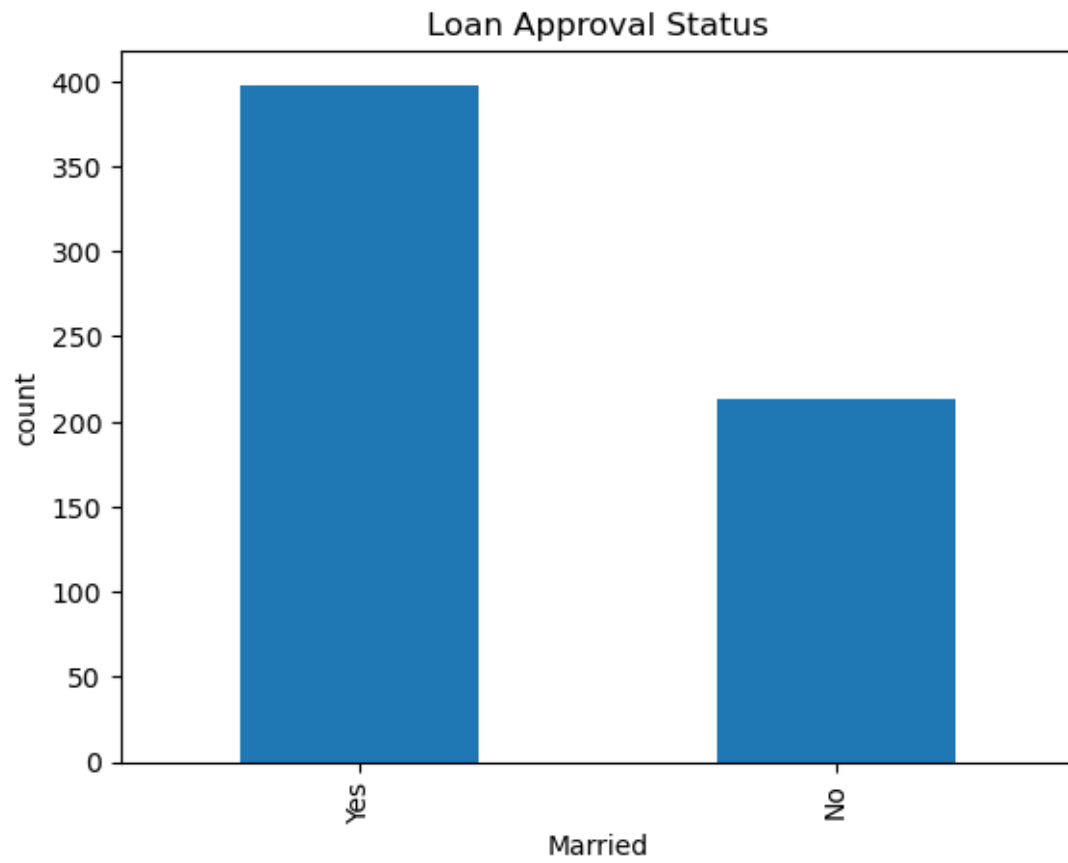
- 'Education' : Comparing loan approvals for individuals with different educational backgrounds (graduate or not) can provide insights into the influence of education on loan approval.
- 'married': Analyzing loan approvals based on marital status can reveal if being married has any impact on the likelihood of loan approval.
- 'dependents' : Exploring loan approvals across different levels of dependents will help understand if the number of dependents affects the loan approval decision.
- 'self employed': Analyzing loan approvals for self-employed individuals versus those who are not self-employed can highlight any differences in loan approval rates based on employment type.
- 'applicantIncome' and 'coapplicantIncome': Plotting the distribution of applicant and co-applicant incomes can reveal any income-related patterns in loan approvals.
- 'loanAmount' and 'loanAmountTerm': Analyzing the distribution of loan amounts and loan amount terms can provide insights into the typical loan sizes and terms that are being approved.
- 'creditHistory': Examining loan approvals based on credit history (with a binary value of 0 or 1) can reveal the impact of creditworthiness on loan approval.
- 'propertyArea': Plotting loan approvals based on the property area (urban, semi-urban, or rural) can highlight any regional differences in loan approval rates.

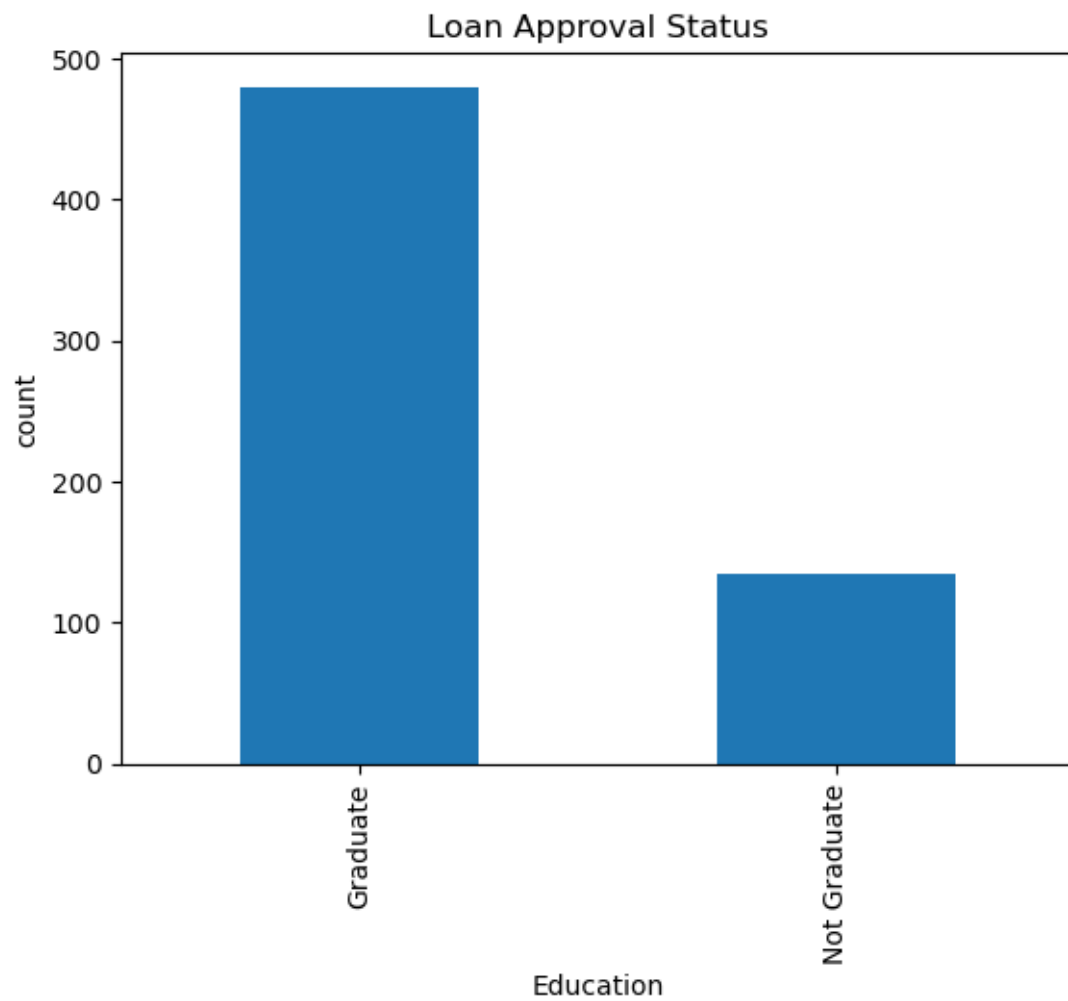
```
[6]: #Sorting Categorical and Numerical columns
categorical = ['Loan_Status', 'Gender', 'Married', 'Education', 'Property_Area', '
↳ 'Self_Employed', 'Dependents', 'Loan_Amount_Term', 'Credit_History']

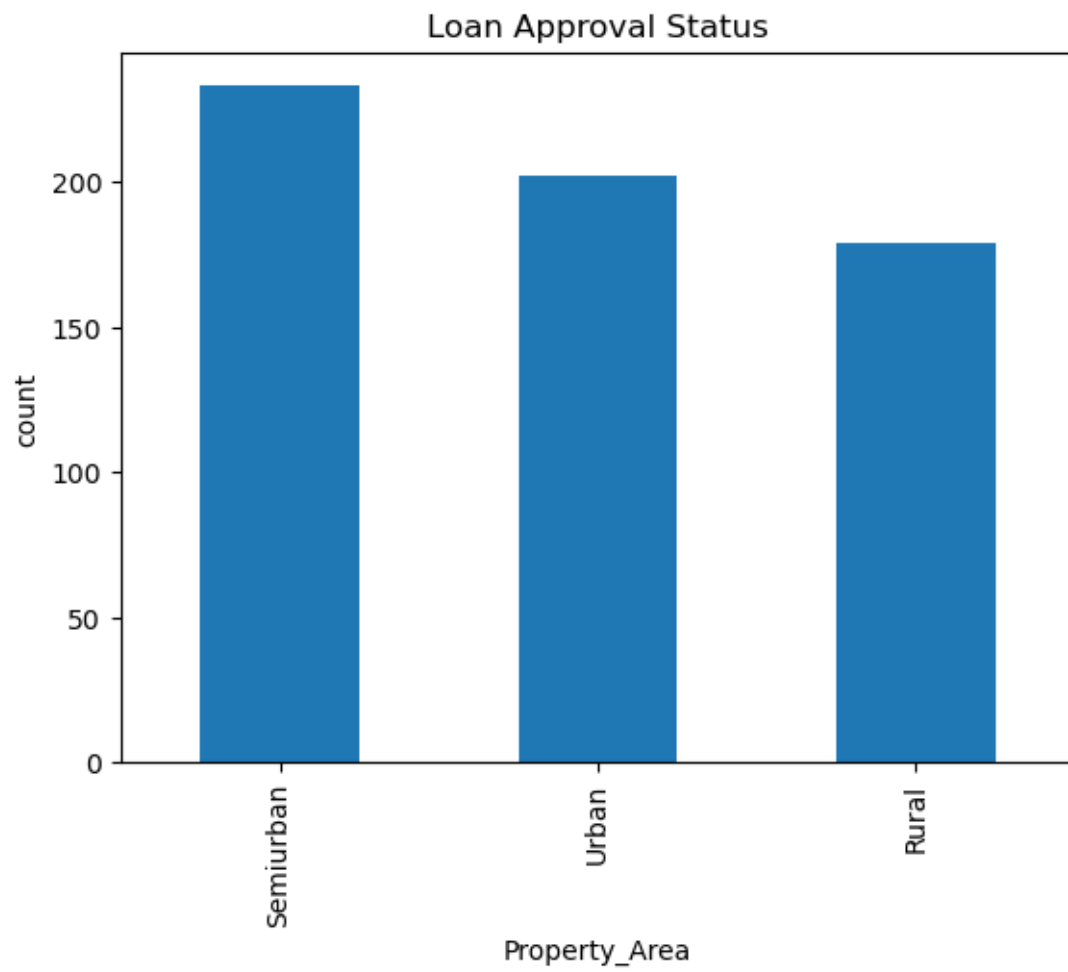
#Plotting all the Categorical columns
for i, column in zip(range(1,10), categorical):
    loan_appr[column].value_counts().plot.bar()
    plt.title('Loan Approval Status')
    plt.xlabel(column)
    plt.ylabel("count")
    plt.show()
```

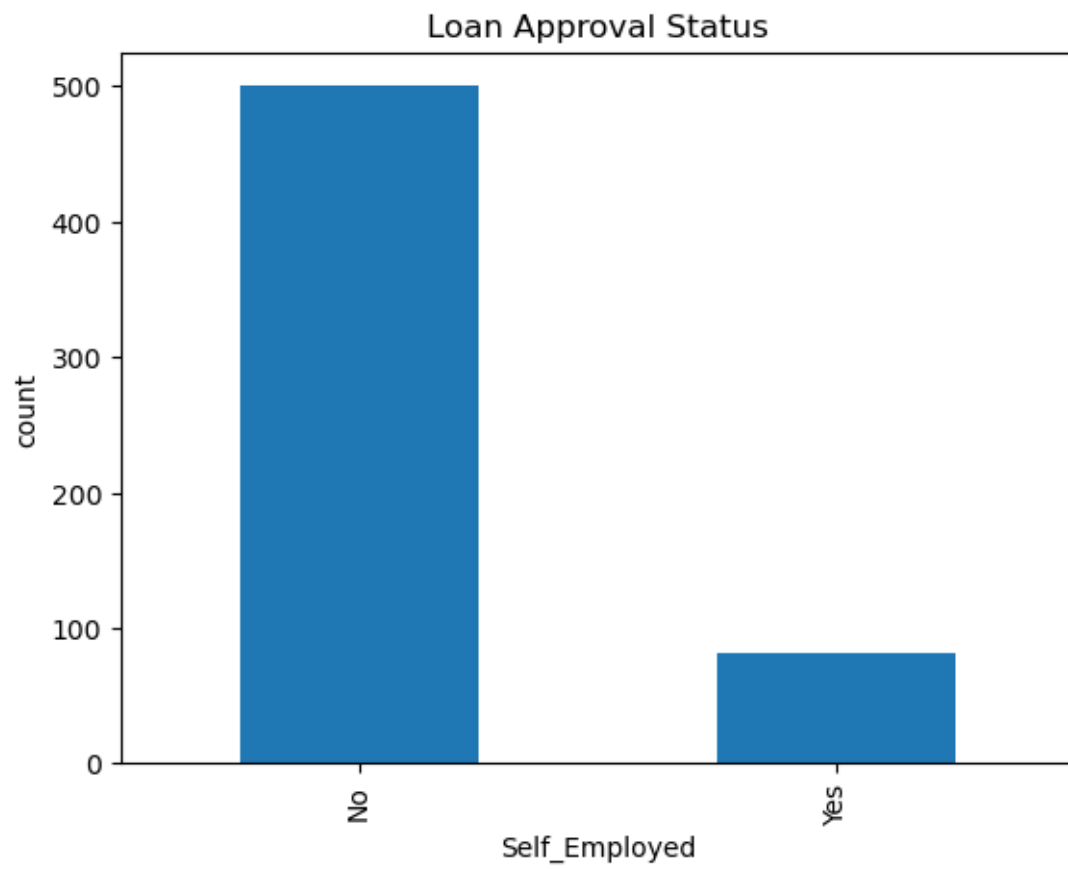


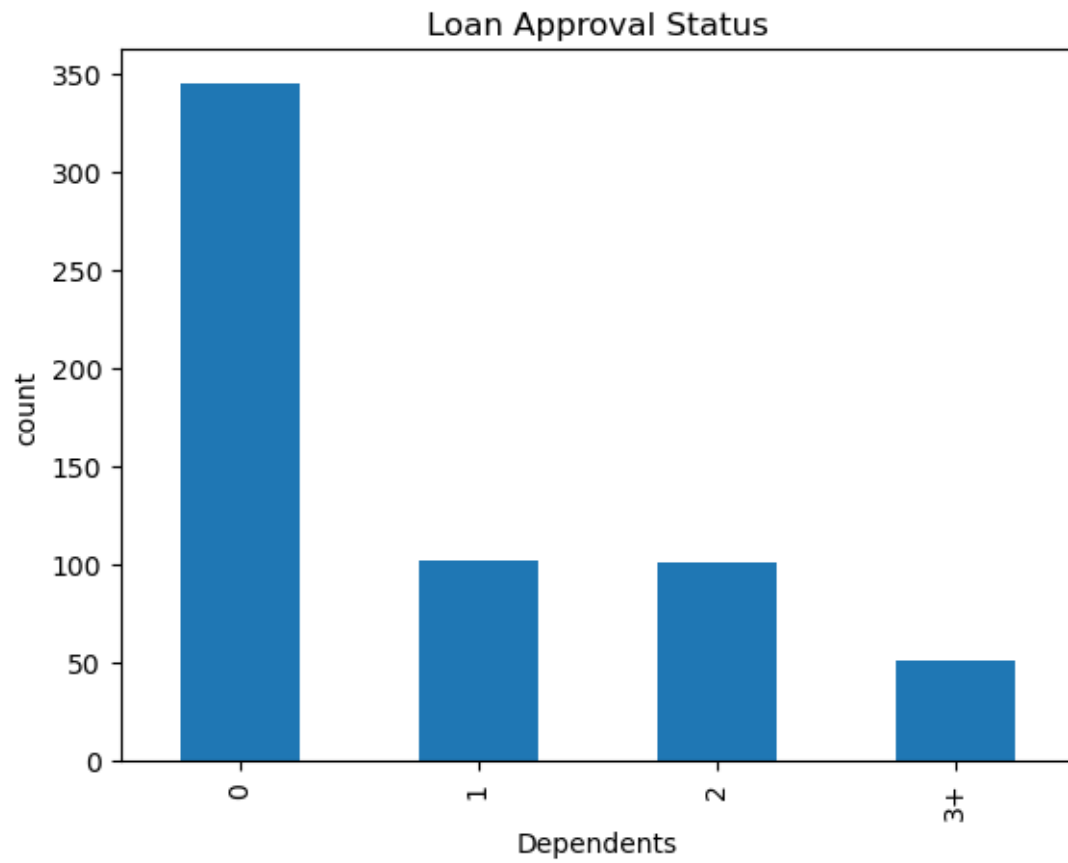


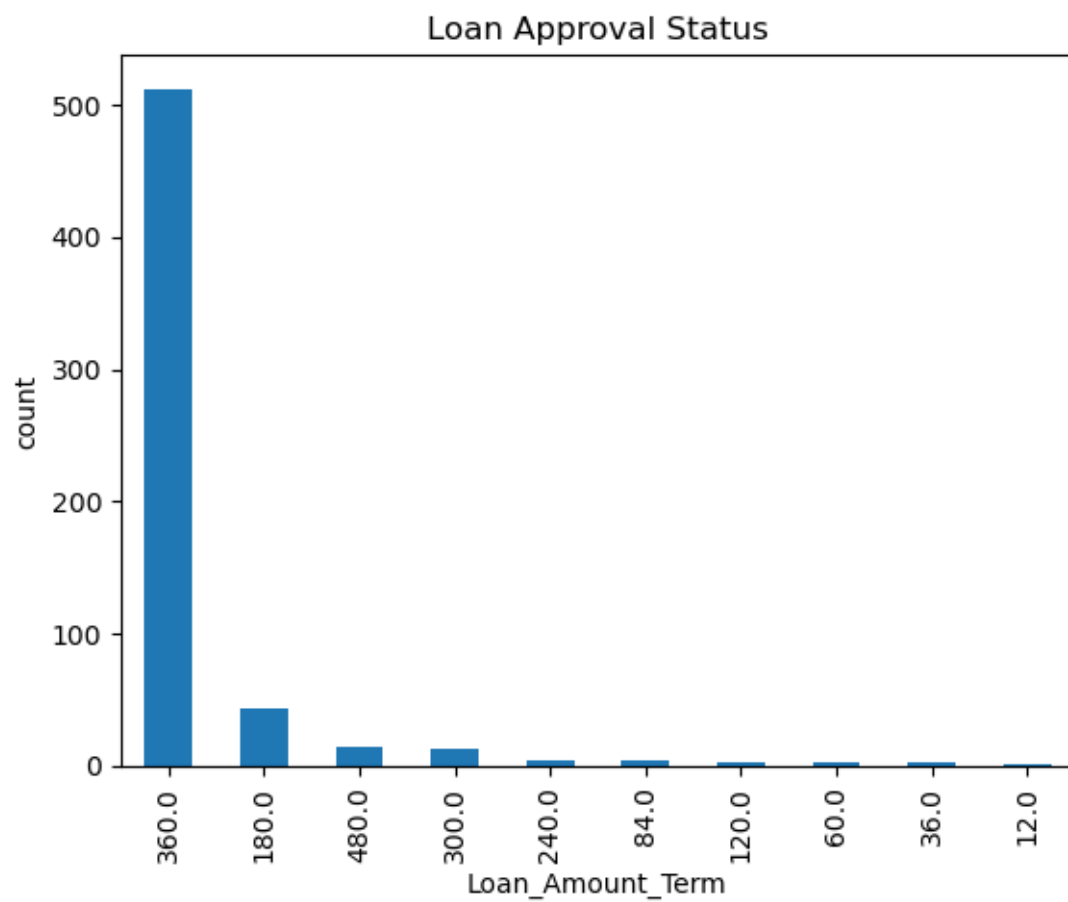















Looking at all the bar plots for the categorical columns, we can make some initial analysis, and observe a trend:

- **Loan_Status:** By the bar plot generated by the **Loan_Status** column we can observe that the loan approval rate is greater compared to the loan rejection rate, we can also further explore the factors influencing loan approval, by comparing with other variables and build a loan approval prediction model
- **Gender :** The higher loan approval rate for males in the **Gender** bar plot indicates a potential gender bias in the loan approval process. we can further analyze other variables and compare with gender column to identify potential factors that contribute to the observed difference in loan approval rates.
- **Married :** The bar plot for the **Married** column indicates higher loan approval rate for the Married marital status, associating being married, with greater financial stability, leading to a positive impact on the loan approval decisions
- **Education:** the bar plot for the **Education** column indicates a higher loan approval rate for educated applicants, suggesting that having a higher level of education has a greater impact on loan approval rate
- **Property Area :** the plot for the **property_area** column showed a almost similar amount


of loans approved for all the semiurban, urban and rural areas, suggesting that the property area is not a major factor in loan approval decision.


- **Self_Employed:** the self employed plot indicated a higher number of loan approvals for the salaried applicants. suggesting a preference for the salaried applicants in the loan approval decisions
 - **Dependents:** the dependents plot, indicates that having a low number or zero dependents, played a huge impact on the loan approval decisions, with a majority of loan approved for applicants with zero dependents.
 - **Credit History:** The Credit history plot indicated that credit history can certainly be considered to be a factor in the loan approval process, with almost 475 out of 564 loans approved for candidates with a credit history
-


1.1 Plotting the Numerical Columns

```
[7]: #Sorting Numerical Columns
Numerical = ['ApplicantIncome', 'CoapplicantIncome', 'LoanAmount']

sns.set_theme()
#Plotting all the numerical columns
for i, column in zip(range(1,4), Numerical):
    sns.histplot(loan_appr[column])
    plt.show()
```







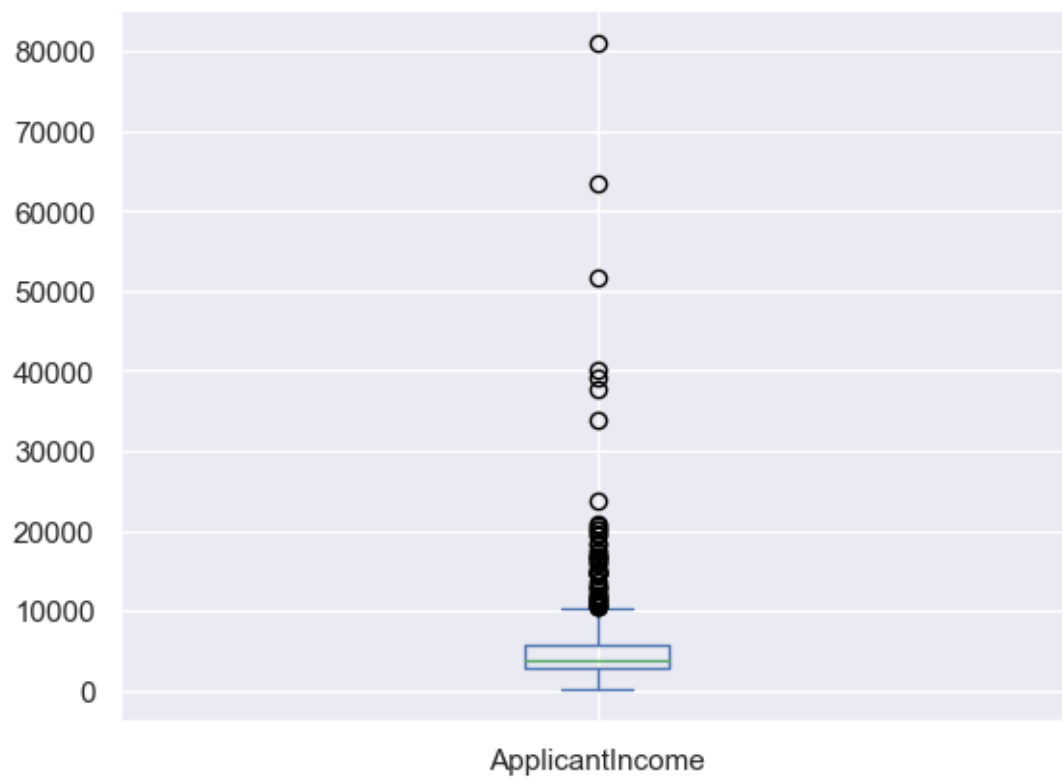
Looking at the plot it can be inferred that the distribution of **Applicant Incomes**, is Right-skewed, with majority of applicants, incomes ranging from 0 - 10000.

We see a similar distribution as that of the applicant income. Majority of **coapplicant's** income ranges from 0 to 5000. We also see a lot of outliers in the coapplicant income and it is not normally distributed.

We see a fairly normal distribution which is slightly right-skewed for **LoanAmount**

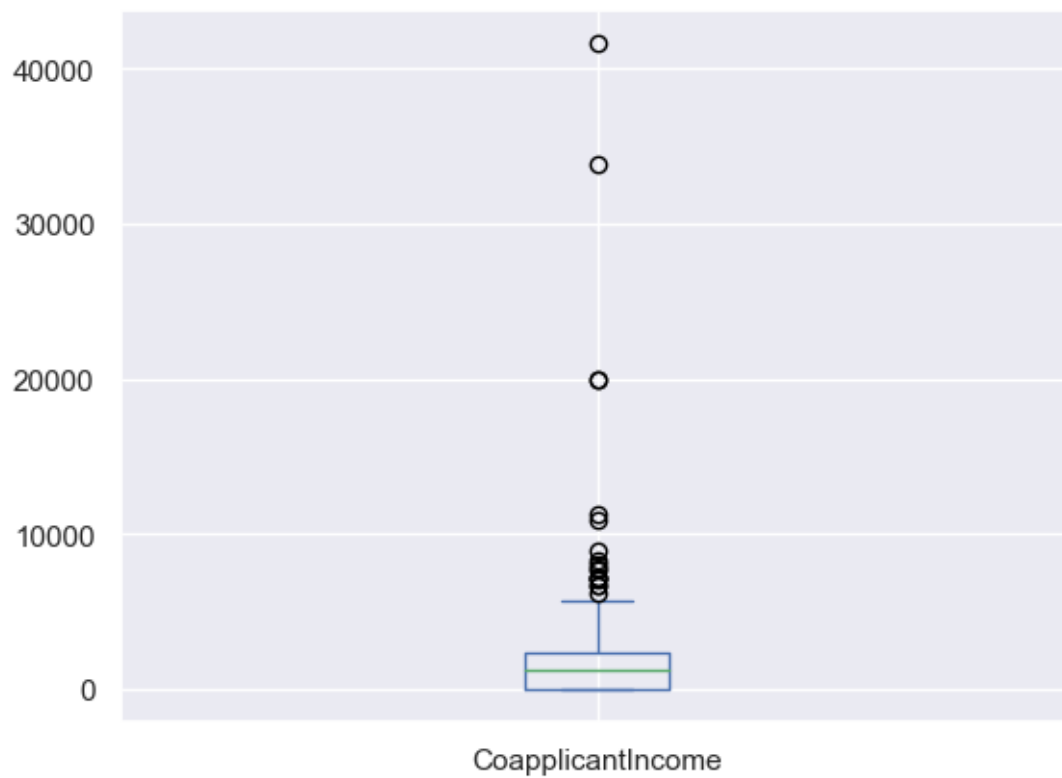
The Numerical variables can further be plotted as boxplots to check for outliers

```
[8]: loan_appr['ApplicantIncome'].plot.box()
plt.show()
```



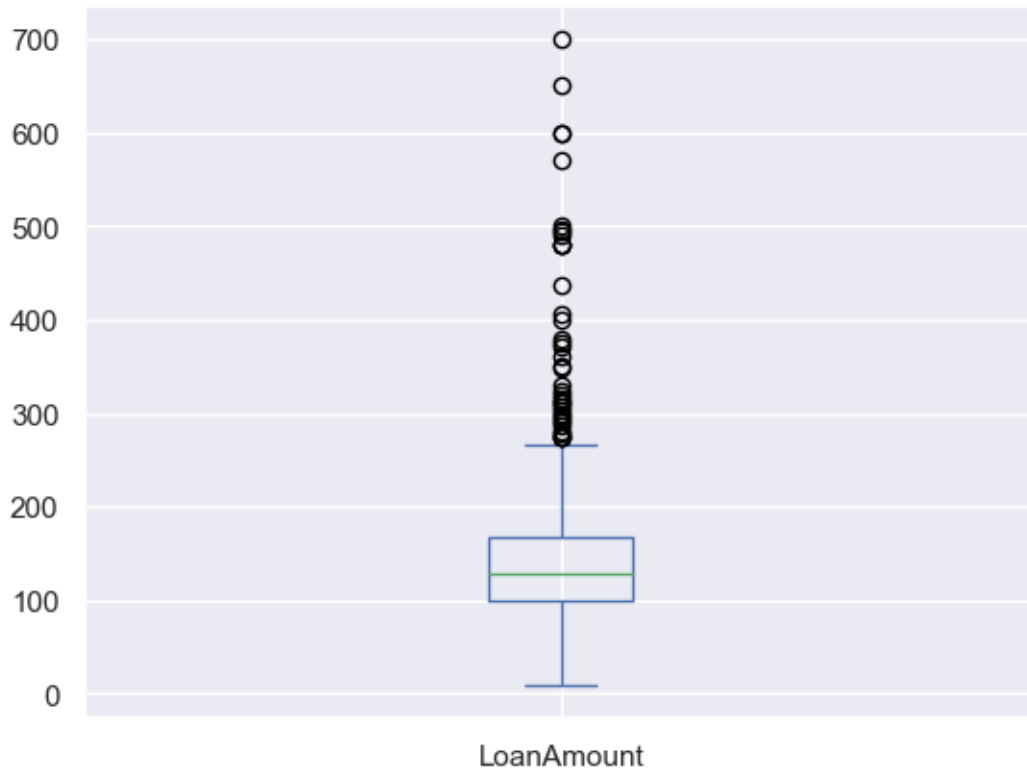
The boxplot shows the presence of a lot of outliers. which could be due to different education level

```
[9]: loan_appr['CoapplicantIncome'].plot.box()  
plt.show()
```



Just like applicant income, the CoapplicantIncome also has a lot of outliers

```
[10]: loan_appr['LoanAmount'].plot.box()  
plt.show()
```



After looking at all the variables individually, let's explore the relationship of these variables with the Loan_Status column, to determine the factors effecting loan approval decisions


1.2 Categorical Variables vs Loan Status

Gender vs Loan Status

```
[11]: pd.crosstab(loan_appr['Gender'], loan_appr['Loan_Status'])
```

```
[11]: Loan_Status    N    Y
Gender
Female           37   75
Male            150  339
```

```
[12]: gender_loan_status = pd.crosstab(loan_appr['Gender'], loan_appr['Loan_Status'])
gender_loan_status.plot(stacked = True)
plt.show()
```




From the bar chart it can be seen that the proportion of loan approval for male and female applicants is more or less same.

Married vs Loan Status

```
[13]: pd.crosstab(loan_appr['Married'], loan_appr['Loan_Status'])
```

```
[13]: Loan_Status    N    Y
Married
No              79  134
Yes            113  285
```

```
[14]: married_loan_status = pd.crosstab(loan_appr['Married'],
↳ loan_appr['Loan_Status'])
married_loan_status.plot.bar(stacked = True, rot = 0)
plt.show()
```

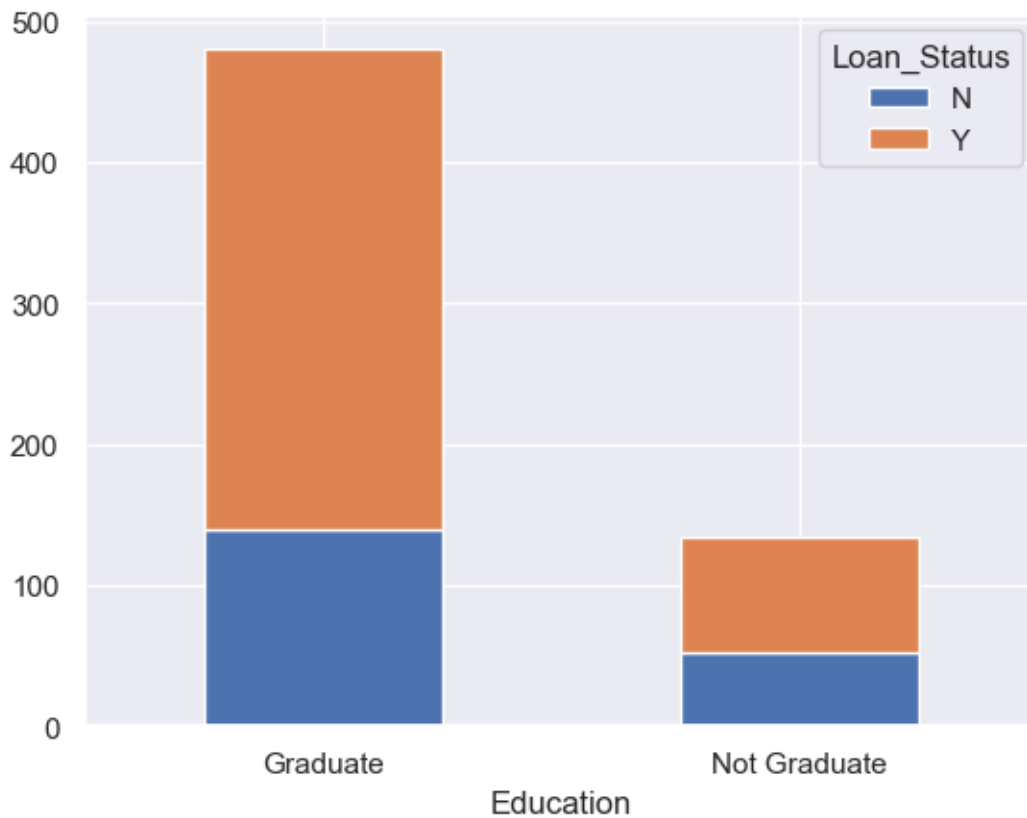


From the bar plot we can see that the proportion of married applicants is higher for the approved loans

```
[15]: pd.crosstab(loan_appr['Education'], loan_appr['Loan_Status'])
```

```
[15]: Loan_Status      N      Y
Education
Graduate           140   340
Not Graduate        52    82
```

```
[16]: education_loan_status = pd.crosstab(loan_appr['Education'],
↳ loan_appr['Loan_Status'])
education_loan_status.plot.bar(stacked = True, rot = 0)
plt.show()
```



we can see that a higher number of graduates received loan approvals compared to non-graduates

```
[17]: pd.crosstab(loan_appr['Property_Area'], loan_appr['Loan_Status'])
```

```
[17]: Loan_Status      N      Y
Property_Area
Rural              69    110
Semiurban          54    179
Urban              69    133
```

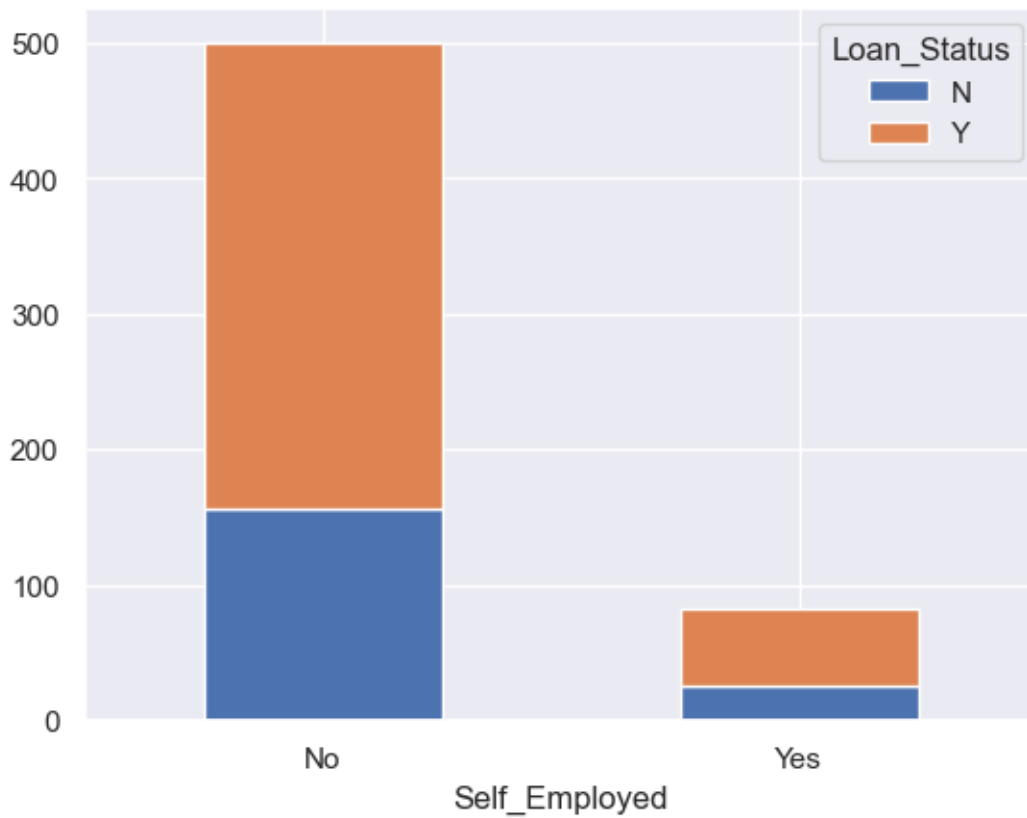
```
[18]: property_loan_status = pd.crosstab(loan_appr['Property_Area'],
    ↪ loan_appr['Loan_Status'])
property_loan_status.plot.bar(stacked = True, rot = 0)
plt.show()
```




```
[19]: pd.crosstab(loan_appr['Self_Employed'], loan_appr['Loan_Status'])
```

```
[19]: Loan_Status      N      Y
Self_Employed
No           157   343
Yes           26    56
```

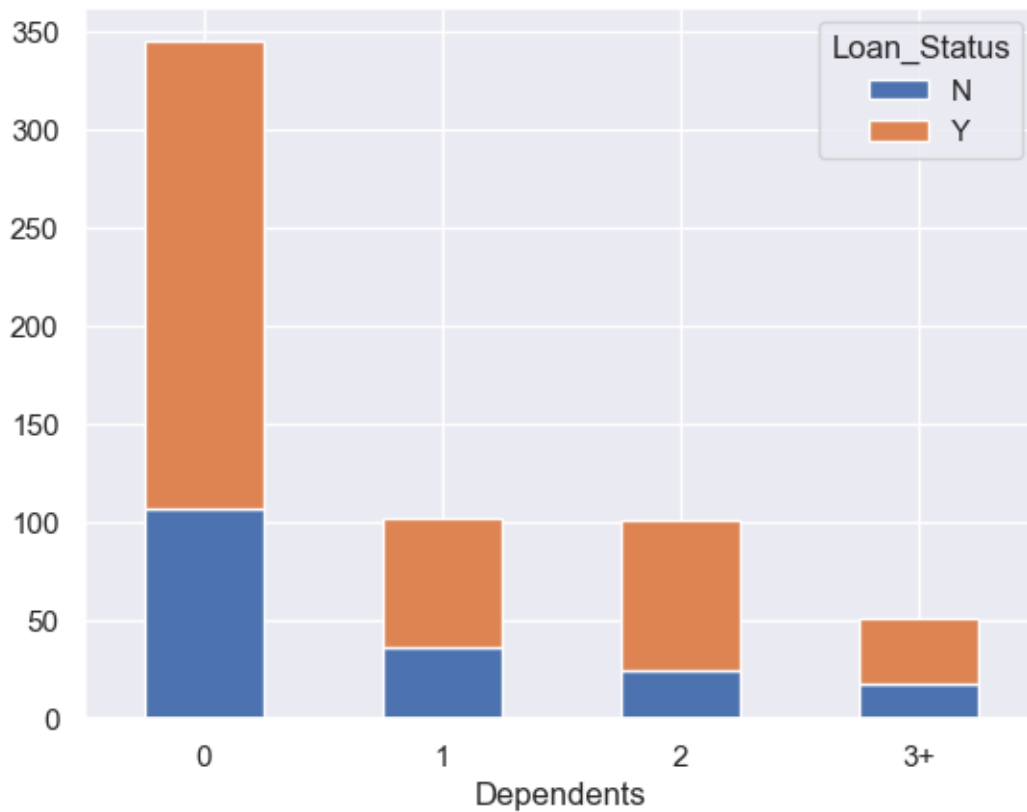
```
[20]: Self_Employed_loan_status = pd.crosstab(loan_appr['Self_Employed'],
↳ loan_appr['Loan_Status'])
Self_Employed_loan_status.plot.bar(stacked = True, rot = 0)
plt.show()
```



```
[21]: pd.crosstab(loan_appr['Dependents'], loan_appr['Loan_Status'])
```

```
[21]: Loan_Status    N    Y
Dependents
0             107  238
1              36   66
2              25   76
3+             18   33
```

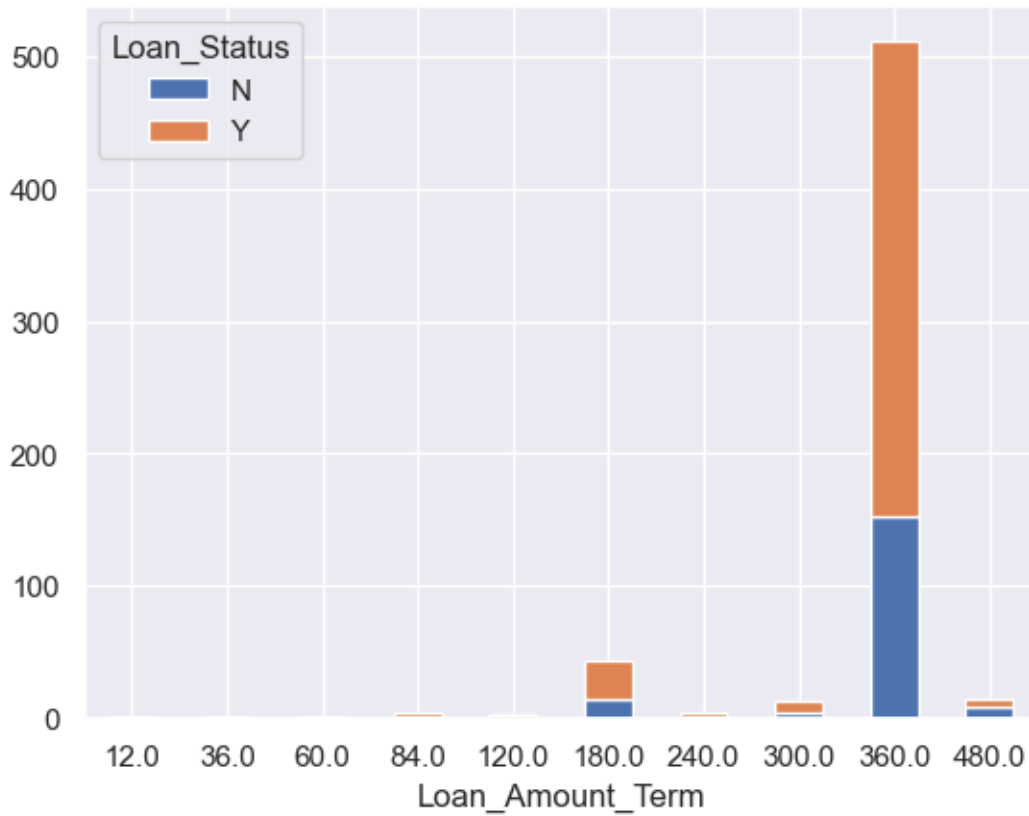
```
[22]: dependents_loan_status = pd.crosstab(loan_appr['Dependents'],
↳ loan_appr['Loan_Status'])
dependents_loan_status.plot.bar(stacked = True, rot = 0)
plt.show()
```



```
[23]: pd.crosstab(loan_appr['Loan_Amount_Term'], loan_appr['Loan_Status'])
```

```
[23]: Loan_Status      N      Y
Loan_Amount_Term
12.0                0      1
36.0                2      0
60.0                0      2
84.0                1      3
120.0               0      3
180.0              15     29
240.0               1      3
300.0               5      8
360.0             153    359
480.0               9      6
```


```
[24]: term_loan_status = pd.crosstab(loan_appr['Loan_Amount_Term'],
↳ loan_appr['Loan_Status'])
term_loan_status.plot.bar(stacked = True, rot = 0)
plt.show()
```



```
[25]: pd.crosstab([loan_appr['Credit_History'], loan_appr['Loan_Status']])
```

```
[25]: Loan_Status      N      Y
Credit_History
0.0              82      7
1.0              97    378
```

```
[26]: credit_history = pd.crosstab([loan_appr['Credit_History'],
↳ loan_appr['Loan_Status']])
credit_history.plot.bar(stacked = True, rot = 0)
plt.show()
```



the bar plot suggests that people with credit history as 1 are more likely to get their loans approved

1.3 Numerical Variables Vs Loan Status

1.3.1 Applicant Income vs Loan Status

```
[27]: pd.crosstab(loan_appr['ApplicantIncome'], loan_appr['Loan_Status'])
```

```
[27]: Loan_Status      N  Y
ApplicantIncome
150          1  0
210          0  1
416          1  0
645          0  1
674          0  1
...          ..  ..
39147         0  1
39999         0  1
51763         0  1
63337         0  1
81000         1  0
```

[505 rows x 2 columns]


```
[28]: print(loan_appr.groupby('Loan_Status')['ApplicantIncome'].mean())
```

```
Loan_Status
N    5446.078125
Y    5384.068720
Name: ApplicantIncome, dtype: float64
```

```
[29]: loan_appr['IncomeRange'] = pd.cut(loan_appr['ApplicantIncome'], bins = [0,
↳2500, 5000, 10000, 20000, 40000, 81000])
income_status = pd.crosstab(loan_appr['IncomeRange'], loan_appr['Loan_Status'])
print(income_status)
income_status.plot.bar(stacked = True, rot = 45)
```

Loan_Status	N	Y
IncomeRange		
(0, 2500]	34	74
(2500, 5000]	99	216
(5000, 10000]	42	97
(10000, 20000]	13	27
(20000, 40000]	3	6
(40000, 81000]	1	2

```
[29]: <Axes: xlabel='IncomeRange'>
```




The bar plot suggests, that the applicant income does not effect the loan approval process as much as anticipated, as we can notice that for every income range the loan approval and rejection is almost similar, suggesting that it is some other factor that determine's the loan approval

But, while applicant income is not a factor in loan approval, let's see if it plays any role in loan amount.

```
[30]: sns.scatterplot(data = loan_appr, x = 'LoanAmount', y = 'ApplicantIncome', hue='Loan_Status', hue_order = ['N', 'Y'])
```

```
[30]: <Axes: xlabel='LoanAmount', ylabel='ApplicantIncome'>
```



The scatterplot suggests, that although the Applicant Income does not have any effect on the Loan Approval, it definitely has an impact on the loan amount showing us a linear positive correlation trend, as seen by the scatterplot, that the higher the income of the applicant, the higher loan amount is approved for the applicant

1.3.2 Applicant Total Income Vs Loan Status

```
[31]: loan_appr['total_income'] = loan_appr['ApplicantIncome'] + \
      ↪ loan_appr['CoapplicantIncome']
print(loan_appr['total_income'].describe())
loan_appr['total_income_range'] = pd.cut(loan_appr['total_income'], bins = [0, \
      ↪ 2500, 5000, 10000, 20000, 40000, 81000])
total_income_status = pd.crosstab(loan_appr['total_income_range'], \
      ↪ loan_appr['Loan_Status'])
print(total_income_status)
```

```
count      614.000000
mean       7024.705081
std        6458.663872
min        1442.000000
25%        4166.000000
50%        5416.500000
```




```

75%      7521.750000
max      81000.000000
Name: total_income, dtype: float64
Loan_Status    N    Y
total_income_range
(0, 2500]      14   10
(2500, 5000]   75  171
(5000, 10000]  74  183
(10000, 20000] 20   49
(20000, 40000]  7    6
(40000, 81000]  2    3

```

```
[32]: total_income_status.plot.bar(stacked = True, rot = 45)
```

```
[32]: <Axes: xlabel='total_income_range'>
```



This graph also does not directly suggest that having a higher income has an effect on loan approval

1.3.3 Loan Amount Vs. Loan Status

```
[33]: pd.crosstab(loan_appr['LoanAmount'], loan_appr['Loan_Status'])
```

```
[33]: Loan_Status  N  Y
LoanAmount
9.0             1  0
17.0            0  1
25.0            0  2
26.0            0  1
30.0            0  2
...           ..  ..
500.0           1  0
570.0           1  0
600.0           0  2
650.0           0  1
700.0           0  1
```

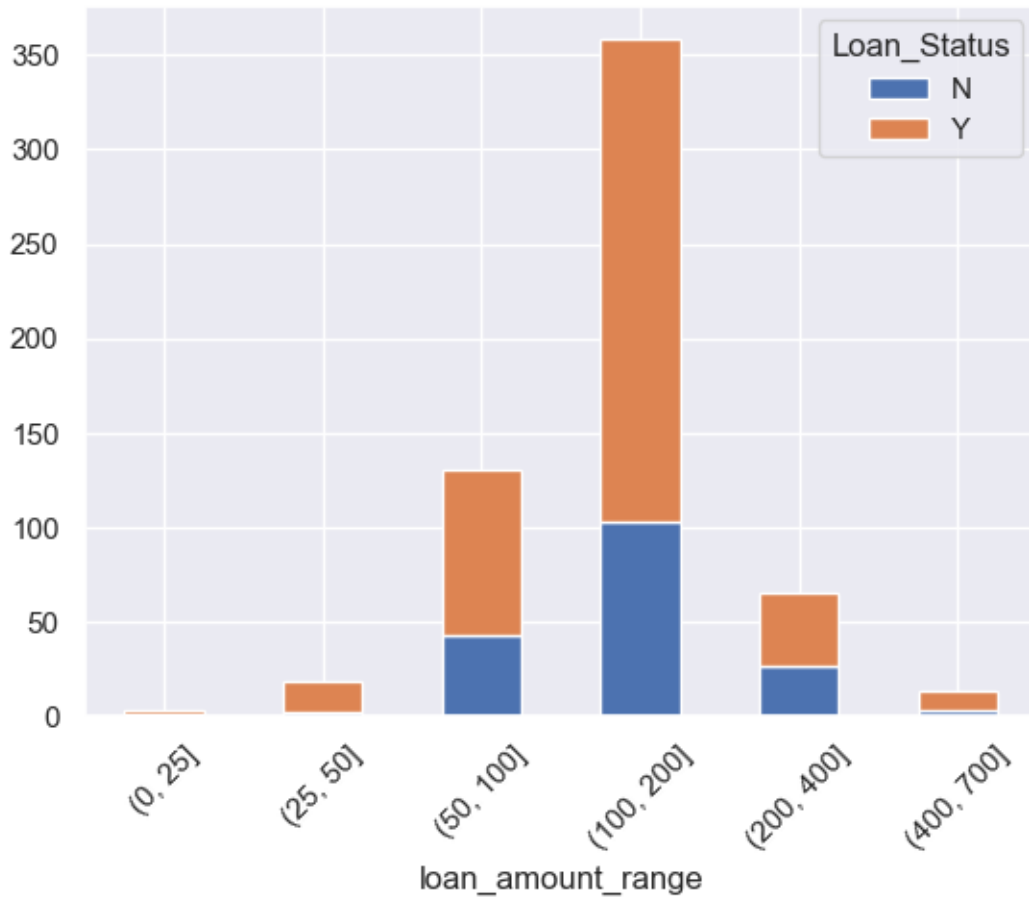
[203 rows x 2 columns]

```
[34]: loan_appr['loan_amount_range'] = pd.cut(loan_appr['LoanAmount'], bins = [0, 25,
↳50, 100, 200, 400, 700])
pd.crosstab(loan_appr['loan_amount_range'], loan_appr['Loan_Status'])
```

```
[34]: Loan_Status      N    Y
loan_amount_range
(0, 25]             1     3
(25, 50]            3    16
(50, 100]           43    88
(100, 200]          103   255
(200, 400]           27    39
(400, 700]           4     10
```

```
[35]: amount_status = pd.crosstab(loan_appr['loan_amount_range'],
↳loan_appr['Loan_Status'])
amount_status.plot.bar(stacked = True, rot = 45)
```

```
[35]: <Axes: xlabel='loan_amount_range'>
```



The plot suggests that the proportion of approval rate for lower and average loan amounts, is higher compared to higher and very high loan amounts

```
[36]: loan_appr = loan_appr.drop(['total_income'], axis = 1)
```

```
[37]: correlation_matrix = loan_appr.corr()
print(correlation_matrix)
```

	ApplicantIncome	CoapplicantIncome	LoanAmount	\
ApplicantIncome	1.000000	-0.116605	0.570909	
CoapplicantIncome	-0.116605	1.000000	0.188619	
LoanAmount	0.570909	0.188619	1.000000	
Loan_Amount_Term	-0.045306	-0.059878	0.039447	
Credit_History	-0.014715	-0.002056	-0.008433	

	Loan_Amount_Term	Credit_History
ApplicantIncome	-0.045306	-0.014715
CoapplicantIncome	-0.059878	-0.002056
LoanAmount	0.039447	-0.008433

Loan_Amount_Term	1.000000	0.001470
Credit_History	0.001470	1.000000

```

/var/folders/xf/djdr51tx47d1mk2ynq3k0hl80000gn/T/ipykernel_30093/3411523848.py:1
: FutureWarning: The default value of numeric_only in DataFrame.corr is
deprecated. In a future version, it will default to False. Select only valid
columns or specify the value of numeric_only to silence this warning.
correlation_matrix = loan_appr.corr()

```

1.4 Cleaning the Data

We will now clean the data, as it is inconsistent and lacking in certain areas, and contains some errors

```
[38]: loan_appr.isnull().sum()
```

```

[38]: Loan_ID          0
      Gender          13
      Married         3
      Dependents      15
      Education        0
      Self_Employed   32
      ApplicantIncome  0
      CoapplicantIncome 0
      LoanAmount      22
      Loan_Amount_Term 14
      Credit_History   50
      Property_Area    0
      Loan_Status      0
      IncomeRange      0
      total_income_range 0
      loan_amount_range 22
      dtype: int64

```

we can observe that there are missing values in Gender, Married, Dependents, Self_Employed, LoanAmount, Loan_Amount_Term and Credit_History features. We will treat the missing values in all the features one by one.

```
[39]: loan_appr_clean = loan_appr.dropna(axis = 0)
```

```
[40]: loan_appr_clean.isnull().sum()
```

```

[40]: Loan_ID          0
      Gender          0
      Married         0
      Dependents      0
      Education        0
      Self_Employed   0

```

```

ApplicantIncome      0
CoapplicantIncome     0
LoanAmount            0
Loan_Amount_Term      0
Credit_History        0
Property_Area         0
Loan_Status           0
IncomeRange           0
total_income_range    0
loan_amount_range     0
dtype: int64

```

```
[41]: loan_appr_clean.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 480 entries, 1 to 613
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Loan_ID               480 non-null   object
1   Gender                480 non-null   object
2   Married               480 non-null   object
3   Dependents            480 non-null   object
4   Education              480 non-null   object
5   Self_Employed         480 non-null   object
6   ApplicantIncome       480 non-null   int64
7   CoapplicantIncome     480 non-null   float64
8   LoanAmount            480 non-null   float64
9   Loan_Amount_Term      480 non-null   float64
10  Credit_History         480 non-null   float64
11  Property_Area          480 non-null   object
12  Loan_Status            480 non-null   object
13  IncomeRange            480 non-null   category
14  total_income_range     480 non-null   category
15  loan_amount_range      480 non-null   category
dtypes: category(3), float64(4), int64(1), object(8)
memory usage: 54.7+ KB

```

We will now clean the column names, to fit the Python snakecase

```

[42]: new_columns = []
      for c in loan_appr_clean.columns:
          new_columns.append(c.lower())

      loan_appr_clean.columns = new_columns
      print(loan_appr_clean.columns)

```

```

Index(['loan_id', 'gender', 'married', 'dependents', 'education',
      'self_employed', 'applicantincome', 'coapplicantincome', 'loanamount',

```

```

        'loan_amount_term', 'credit_history', 'property_area', 'loan_status',
        'incomerange', 'total_income_range', 'loan_amount_range'],
dtype='object')

```

```

[43]: loan_appr_clean.columns = ['loan_id', 'gender', 'married', 'dependents', 'education',
        'self_employed', 'applicant_income', 'coapplicant_income', 'loan_amount',
        'loan_amount_term', 'credit_history', 'property_area', 'loan_status',
        'income_range', 'total_income_range', 'loan_amount_range']
print(loan_appr_clean.columns)

```

```

Index(['loan_id', 'gender', 'married', 'dependents', 'education',
       'self_employed', 'applicant_income', 'coapplicant_income',
       'loan_amount', 'loan_amount_term', 'credit_history', 'property_area',
       'loan_status', 'income_range', 'total_income_range',
       'loan_amount_range'],
      dtype='object')

```

```

[51]: loan_appr_clean.describe(include = 'all')

```

```

[51]:
      loan_id  gender  married  dependents  education  self_employed \
count      480    480      480         480         480         480
unique      480      2        2          4          2          2
top    LP001003  Male      Yes          0  Graduate          No
freq          1    394      311        274        383        414
mean         NaN     NaN     NaN         NaN         NaN         NaN
std          NaN     NaN     NaN         NaN         NaN         NaN
min          NaN     NaN     NaN         NaN         NaN         NaN
25%          NaN     NaN     NaN         NaN         NaN         NaN
50%          NaN     NaN     NaN         NaN         NaN         NaN
75%          NaN     NaN     NaN         NaN         NaN         NaN
max          NaN     NaN     NaN         NaN         NaN         NaN

      applicant_income  coapplicant_income  loan_amount  loan_amount_term \
count      480.000000      480.000000  480.000000      480.000000
unique          NaN          NaN          NaN          NaN
top          NaN          NaN          NaN          NaN
freq          NaN          NaN          NaN          NaN
mean      5364.231250      1581.093583  144.735417      342.050000
std      5668.251251      2617.692267   80.508164      65.212401
min       150.000000       0.000000    9.000000      36.000000
25%      2898.750000       0.000000  100.000000      360.000000
50%      3859.000000      1084.500000  128.000000      360.000000
75%      5852.500000      2253.250000  170.000000      360.000000
max      81000.000000     33837.000000  600.000000      480.000000

      credit_history  property_area  loan_status  income_range \
count      480.000000          480          480          480

```

unique	NaN	3	2	6
top	NaN	Semiurban	Y	(2500, 5000]
freq	NaN	191	332	245
mean	0.854167	NaN	NaN	NaN
std	0.353307	NaN	NaN	NaN
min	0.000000	NaN	NaN	NaN
25%	1.000000	NaN	NaN	NaN
50%	1.000000	NaN	NaN	NaN
75%	1.000000	NaN	NaN	NaN
max	1.000000	NaN	NaN	NaN

	total_income_range	loan_amount_range	loan_amount_log
count	480	480	480.000000
unique	6	6	NaN
top	(5000, 10000]	(100, 200]	NaN
freq	198	290	NaN
mean	NaN	NaN	4.848336
std	NaN	NaN	0.510329
min	NaN	NaN	2.197225
25%	NaN	NaN	4.605170
50%	NaN	NaN	4.852030
75%	NaN	NaN	5.135798
max	NaN	NaN	6.396930

```
[54]: print(loan_appr_clean.head())
      print(loan_appr_clean.tail())
```

	loan_id	gender	married	dependents	education	self_employed	\
1	LP001003	Male	Yes	1	Graduate	No	
2	LP001005	Male	Yes	0	Graduate	Yes	
3	LP001006	Male	Yes	0	Not Graduate	No	
4	LP001008	Male	No	0	Graduate	No	
5	LP001011	Male	Yes	2	Graduate	Yes	

	applicant_income	coapplicant_income	loan_amount	loan_amount_term	\
1	4583	1508.0	128.0	360.0	
2	3000	0.0	66.0	360.0	
3	2583	2358.0	120.0	360.0	
4	6000	0.0	141.0	360.0	
5	5417	4196.0	267.0	360.0	

	credit_history	property_area	loan_status	income_range	total_income_range	\
1	1.0	Rural	N	(2500, 5000]	(5000, 10000]	
2	1.0	Urban	Y	(2500, 5000]	(2500, 5000]	
3	1.0	Urban	Y	(2500, 5000]	(2500, 5000]	
4	1.0	Urban	Y	(5000, 10000]	(5000, 10000]	
5	1.0	Urban	Y	(5000, 10000]	(5000, 10000]	

	loan_amount_range	loan_amount_log
1	(100, 200]	4.852030
2	(50, 100]	4.189655
3	(100, 200]	4.787492
4	(100, 200]	4.948760
5	(200, 400]	5.587249

	loan_id	gender	married	dependents	education	self_employed	\
609	LP002978	Female	No	0	Graduate	No	
610	LP002979	Male	Yes	3+	Graduate	No	
611	LP002983	Male	Yes	1	Graduate	No	
612	LP002984	Male	Yes	2	Graduate	No	
613	LP002990	Female	No	0	Graduate	Yes	

	applicant_income	coapplicant_income	loan_amount	loan_amount_term	\
609	2900	0.0	71.0	360.0	
610	4106	0.0	40.0	180.0	
611	8072	240.0	253.0	360.0	
612	7583	0.0	187.0	360.0	
613	4583	0.0	133.0	360.0	

	credit_history	property_area	loan_status	income_range	\
609	1.0	Rural	Y	(2500, 5000]	
610	1.0	Rural	Y	(2500, 5000]	
611	1.0	Urban	Y	(5000, 10000]	
612	1.0	Urban	Y	(5000, 10000]	
613	0.0	Semiurban	N	(2500, 5000]	


	total_income_range	loan_amount_range	loan_amount_log
609	(2500, 5000]	(50, 100]	4.262680
610	(2500, 5000]	(25, 50]	3.688879
611	(5000, 10000]	(200, 400]	5.533389
612	(5000, 10000]	(100, 200]	5.231109
613	(2500, 5000]	(100, 200]	4.890349


1.5 Cleaning Outliers


we saw earlier that some of the variables in the dataset contains outliers. due to these outliers our variables are right skewed, which is having a significant effect on the mean and standard deviation affecting our distribution.

```
[44]: Numerical_clean = ['applicant_income', 'coapplicant_income', 'loan_amount']

sns.set_theme()
#Plotting all the numerical columns
for i, column in zip(range(1,4), Numerical_clean):
    sns.histplot(loan_appr_clean[column])
    plt.show()
```






To remove this skewness, one method we can use is the log transformation


```
[65]: loan_appr_clean.loc['loan_amount_log'] = np.log(loan_appr_clean['loan_amount'])
loan_appr_clean.loc['applincome_log'] = np.
    ↳log(loan_appr_clean['applicant_income'])
loan_appr_clean.loc['coapplincome_log'] = np.
    ↳log(loan_appr_clean['coapplicant_income'])
```


/Users/chinmay/anaconda3/lib/python3.10/site-packages/pandas/core/arraylike.py:402: RuntimeWarning: divide by zero encountered in log

```
    result = getattr(ufunc, method)(*inputs, **kwargs)
```

```
[67]: sns.histplot(loan_appr_clean['applincome_log'])
plt.show()
sns.histplot(loan_appr_clean['coapplincome_log'])
plt.show()
sns.histplot(loan_appr_clean['loan_amount_log'])
plt.show()
```







1.6 Conclusion

In this project, we conducted a comprehensive analysis of the **Loan+Approval+Prediction** dataset, aiming to gain insights into the factors influencing loan approval and explore potential predictions based on the analysis.

During the **Univariate Analysis** we examined each column/variable individually to understand their distributions and characteristics. The analysis revealed that the majority of applicants which has their loans approved had a **credit history**, were **male**, **educated**, **employed** and **married**, with varying range of income levels.

We further explored the relationships of the variables, and hypothesis we generated in the **Univariate Analysis** in the **Bivariate Analysis**.

Moving on to the **Bivariate Analysis** we explored the relationship of different variables vs the **Loan Approval Status**. Notably we discovered that although we would generally expect a **higher income** to be a important factor in **loan approval**, the research suggests that it is not the case and that the **loan approval rate** was almost similar for **all income ranges**, but the **applicant income** did play a role in the **loan amount approved** - *the higher the income the higher the loan amount approved*. Additionally we also observed a positive correlation between **credit history**

and loan approval, indicating that a having a credit history significantly influenced the loan decision

1.7 Final Predictions

Based on the insights gained from our analysis, we can make the following predictions: - **Income and Loan Approval:** even though the research does not directly support the hypothesis, that a higher income has an impact on the decision making process, we cannot disregard the human bias, a higher income gives a sense of security that the applicant in the future can return the loan. which is supported by the analysis we observed that a higher income approved a higher loan amount for the candidate.

- **Credit History and Loan Approval:** A positive credit history strongly influences the loan approval decision. Applicants with a credit history are more likely to secure a loan.
- **Marrital status, education and no.of dependents:** While not explicitly analyzed, during the univariate analysis, having a marrital status, showed a potential impact on loan approval, also being educated and having less no. of dependents definately are a factor in the loan approval decision.

1.7.1 In conclusion,

while this analysis provides valuable insights into specific factors influencing loan approval, it is important to acknowledge the limitations and the need for further research. A comprehensive prediction requires consideration of additional factors, such as human bias and other relevant variables. Moreover, leveraging advanced machine learning models can improve the accuracy of loan approval predictions.